

1 Subjects

- Bagging
- Boosting

2 Notes

2.1 Decision trees

Decision trees simply split the feature space into a set of rectangles, and then fit a simple model (e.g. a constant) in each one.

So one might say that, we take some input x and ask a series of questions. Is $x < 21$? Then go down the left path of the tree of questions. So decision trees are very simple to understand but gain are also very expressive. One of the major benefits of decision trees, is that they are very easy to read and understand. It's very easy to look at it and figure out what questions it asks and how much does it weigh that question (e.g. what pixels does it look at to understand what the image is, what pixels are "important").

If we just look at a very simple form of decision trees, i.e. binary trees then things will be a bit simpler. Then every question is of the form "if this then go left, otherwise go right" so every node in the tree is a single dividing line in the feature space. If we look at tree where the values of the regions are constant, then we can predict the value of some input x as follows:

$$h(x) = \sum_{\text{regions } r} 1_{[x \in r]} \cdot c_r(x)$$

That is the constant value of the region that x is in.

2.1.1 Growing a regression tree

Using the binary regression tree described earlier, with constant region values, we will then describe how to grow a regression tree (learn).

Given some input of N observations (x_i, y_i) for $i = 1, 2, \dots, N$ with $x_i = (x_{i1}, x_{i2}, \dots, x_{id})$ and y_i being the label or the "true" value, we then need to figure out how to approximate y for some unknown x outside of the training set. The algorithm then needs to decide on the splitting variables and split points, as well as the topology of the tree.

Suppose we have a partition into M regions R_1, R_2, \dots, R_M and we model the response as mentioned earlier:

$$h(x) = \sum_{\text{regions } r} 1_{[x \in r]} \cdot c_r(x)$$

As usual with regression, we can use the squared error measure: $(h(x_i) - y_i)^2$ to evaluate our performance. Then, we can easily see that the best constant for

region R_m is simply the average of y_i which ended up in region R_m , because least squares measure punishes distance the target. I.e. one point which is off by 2 is punished more than two points that are off by 1, and thus the mean is the smallest distance on average:

$$\hat{c}_m = \frac{1}{|D|} \sum_{(x,y) \in D} 1_{[x \in r]} \cdot y$$

Now, actually finding the best partition of the feature space into the R_m optimal regions is, in general, computationally infeasible. Thus, we will proceed with a greedy approximation algorithm.

Consider some splitting variable j and splitting point s , we can then define the pair of half-planes:

$$R_1(j, s) = \{X | X_j \leq s\} \text{ and } R_2(j, s) = \{X | X_j > s\}$$

Note here that the point X can be any point in the feature space and is not necessarily a point from D .

We then seek to find the splitting variable j and split point s that solve:

$$\min_{j,s} \left[\min_{c_1} \left[\sum_{x_i \in R_1(j,s)} (y_i - c_1)^2 \right] + \min_{c_2} \left[\sum_{x_i \in R_2(j,s)} (y_i - c_2)^2 \right] \right]$$

If we use the mean, as mentioned earlier, we can simplify this to:

$$\min_{j,s} \left[\left(\sum_{x_i \in R_1(j,s)} (y_i - \hat{c}_1)^2 \right) + \left(\sum_{x_i \in R_2(j,s)} (y_i - \hat{c}_2)^2 \right) \right]$$

Or, stated differently:

$$\min_{j,s} \left[\left(\sum_{(x,y) \in D, x_j \leq s} (y - \hat{c}_1)^2 \right) + \left(\sum_{(x,y) \in D, x_j > s} (y - \hat{c}_2)^2 \right) \right]$$

For each splitting variable j , we can simply pick s as the midpoints between two x_j values, which results in $|D| - 1$ different split-points to consider.

Now that we know how to compute the “best” split, we simply follow the following greedy algorithm:

- Select best “split” variable j , and the split-point s
- Make a new node in the tree with (j, s) (i.e. a new split)
- Make a child for each new region (outcome of testing j)
- Split the training examples up between the two regions
- Call recursively on the new children (the new regions)
- Stop when done

2.1.2 Size of the tree

How large should the tree become? If we make it too large, we might overfit. Too small and we might not capture the structure of the data properly.

Tree size is a hyper-parameter, and it should be chosen adaptively based on the data. We could, for example, stop if the decrease in error drops below some threshold. This is fairly short-sighted however, as a bad split at some level might result in a crucial split later on.

The usual strategy is to grow a large tree T_0 and stop it when the size of the nodes (i.e. how many data-points we have for each node) drops below some set threshold (e.g. 5) and then pruning this tree.

One way to prune the tree, is to simply compute the accuracy gained by removing on some validation set and greedily deleting the splits that increases accuracy the most. Repeating this as long as it helps.

Limiting the size of the tree is the same as regularizing for decision trees. So variance will go down and bias will go up etc.

2.2 Classification trees

For regression trees, we used the squared error impurity measure $Q_m(T)$:

$$Q_m(T) = \frac{1}{N_m} \sum_{x_i \in R_m} (y_i - \hat{c}_m)^2$$

Where $N_m = |\{x_i | x_i \in R_m\}|$ If the target is a classification outcome, taking values $1, 2, \dots, K$, then this measure is not suitable.

If we let \hat{p}_{mk} be the proportion of observations that had class k in node m . We can then predict the class of points in region m to have class $k(m) = \arg \max_k \hat{p}_{mk}$, i.e. the majority class in region m . Then, different measures of $Q_m(T)$ of node impurity includes the following:

Misclassification error (0-1 Loss)

$$\frac{1}{N_m} \sum_{(x,y) \in R_m} 1_{[k(m) \neq y]} = 1 - \hat{p}_{mk(m)}$$

Gini index

$$\sum_{k \neq k'} \hat{p}_{mk} \hat{p}_{mk'} = \sum_{k=1}^K \hat{p}_{mk} (1 - \hat{p}_{mk})$$

Cross-entropy

$$-\sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$$

2.2.1 Why binary splits?

We can represent multiway splits by a series of binary splits (e.g. a tertiary split can be represented by two binary splits) so we don't lose generality. Furthermore, multiway splits split the data too quickly, so we are left with too little information at the lower levels.

2.3 Ensemble methods

Decision trees have a high variance, just one data-point can completely change the outcome of the tree. When we did the pruning and size-limitation we sought to decrease the variance (through regularization). Another option for reducing variance is ensemble methods.

2.3.1 Bootstrapping datasets

Bootstrapping is a general tool for assessing statistical accuracy. Suppose we have our data-set D as before. The basic idea is then to pick random points from this such that we have B different data-sets D_i that each have a random selection of points from D (which may be overlapping). We can use these data-sets D_i as “independent” data-sets. This method is often used to measure statistics like variance, confidence bounds, std.err. etc. in other, more statistical oriented, domains.

2.3.2 Bagging

Bagging uses bootstrapped data-sets in order to improve on our predictions. What we do in bagging (bootstrap aggregation) is simply to get our B bootstrapped data-sets D_i , train our model on all the different data-sets producing h_1, \dots, h_B (our ensemble of models). We can then perform predictions by doing a majority vote (for classification) or returning the mean of the predictions.

$$h_{\text{bag}}(x) = \frac{1}{B} \sum_{i=1}^B h_i(x)$$